

Introduction to AVX Vectorization using C

DR. MATTHEW SMITH, MSMITH@ASTRO.SWIN.EDU.AU

Simplest Possible Beginnings

Consider this standard C code performing a simple vector computation:

$$c_i = \sqrt{a_i^2 + b_i^2}$$

for $i = 0$ to $N-1$ (Since we are using C)

In this code we use a for loop to go over all N values of a and b to compute C .

Is this code the best possible code?

```
/* Sequential demonstration prepared for AVX Cookies 'n' Code Session
   Dr. Matthew Smith, msmith@astro.swin.edu.au */

#include <stdio.h>
#include <malloc.h>
#include <math.h>

void Compute_C(float *x, float *y, float *z, int N);

int main() {
    int N = 1028;      // N = size of problem, a power of 2 is good.
    float *a, *b, *c; // Three vectors containing N elements
    float error = 0.0; // Used later for testing
    size_t size = N*sizeof(float);
    int i;
    // Allocate memory
    a = (float*)malloc(size); b = (float*)malloc(size); c = (float*)malloc(size);
    // Initialise a and b vectors
    for (i = 0; i < N; i++) {
        a[i] = (float)i; b[i] = (float)(2*i);
    }

    Compute_C(a, b, c, N);

    // Check the result
    for (i = 0; i < N; i++) {
        error = error + c[i] - sqrtf(a[i]*a[i] + b[i]*b[i]);
    }
    printf("Error = %g\n", error);

    free(a); free(b); free(c);
    printf("Computation Complete\n");
    return 0;
}

void Compute_C(float *x, float *y, float *z, int N) {
    int i;
    for (i = 0; i < N; i++) {
        z[i] = sqrtf(x[i]*x[i] + y[i]*y[i]);
    }
}
```

Simplest Possible Beginnings

The makefile here:

```
all:
    g++ main.cpp -c -O3
    g++ main.o -O3 -o test.run
```

We might use `-O3` to get the fastest possible performance. This level of optimization should include vectorization. Perhaps we better check:

Disassemble the object to view the assembly-like code: `objdump -d main.o > outputfile.txt`

Simplest Possible Beginnings

What exactly is going on?

- The computation is happening in the XMM1 and XMM0 registers.
- These computations are sequential: how do we know?

mulss, addss, sqrtss → Sequential operations.

Immediately we can conclude that this code is not making the most out of the hardware.


```
0000000000000000 <_Z9Compute_CPfS_S_i>:
 0: 85 c9          test    %ecx,%ecx
 2: 7e 4f          jle     53 <_Z9Compute_CPfS_S_i+0x53>
 4: 55             push    %rbp
 5: 8d 41 ff       lea     -0x1(%rcx),%eax
 8: 53             push    %rbx
 9: 48 8d 2c 85 04 00 00 lea     0x4(,%rax,4),%rbp
10: 00
11: 31 db          xor     %ebx,%ebx
13: 48 83 ec 28     sub     $0x28,%rsp
17: 66 0f 1f 84 00 00 00 nopw    0x0(%rax,%rax,1)
1e: 00 00
20: f3 0f 10 0c 1f movss   (%rdi,%rbx,1),%xmm1
25: f3 0f 10 04 1e movss   (%rsi,%rbx,1),%xmm0
2a: f3 0f 59 c9     mulss   %xmm1,%xmm1
2e: f3 0f 59 c0     mulss   %xmm0,%xmm0
32: f3 0f 58 c8     addss   %xmm0,%xmm1
36: f3 0f 51 c1     sqrtss  %xmm1,%xmm0
3a: 0f 2e c0        ucomiss %xmm0,%xmm0
3d: 7a 16          jp      55 <_Z9Compute_CPfS_S_i+0x55>
3f: f3 0f 11 04 1a movss   %xmm0, (%rdx,%rbx,1)
44: 48 83 c3 04     add     $0x4,%rbx
48: 48 39 eb        cmp     %rbp,%rbx
4b: 75 d3          jne     20 <_Z9Compute_CPfS_S_i+0x20>
4d: 48 83 c4 28     add     $0x28,%rsp
51: 5b             pop     %rbx
52: 5d             pop     %rbp
53: f3 c3          repz    retq
55: 0f 28 c1        movaps  %xmm1,%xmm0
58: 48 89 54 24 18 mov     %rdx,0x18(%rsp)
5d: 48 89 74 24 10 mov     %rsi,0x10(%rsp)
62: 48 89 7c 24 08 mov     %rdi,0x8(%rsp)
67: e8 00 00 00 00 callq   6c <_Z9Compute_CPfS_S_i+0x6c>
6c: 48 8b 54 24 18 mov     0x18(%rsp),%rdx
71: 48 8b 74 24 10 mov     0x10(%rsp),%rsi
76: 48 8b 7c 24 08 mov     0x8(%rsp),%rdi
7b: eb c2          jmp     3f <_Z9Compute_CPfS_S_i+0x3f>
```

SIMD REGISTERS

SIMD REGISTERS

Computer technology has evolved rapidly over a short period of time.


The complexity of the applications, driven by consumers, has also grown rapidly.



SIMD REGISTERS

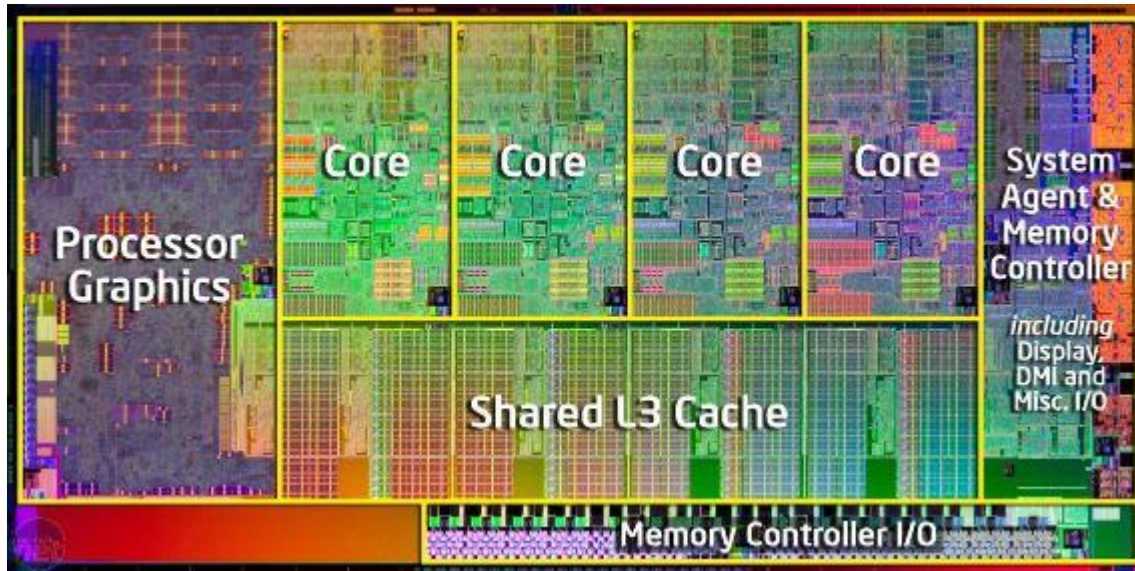
However, in recent years, the clock speed has stagnated – very high frequencies lead to high temperatures.

So what have CPU manufacturers been doing to increase CPU performance?

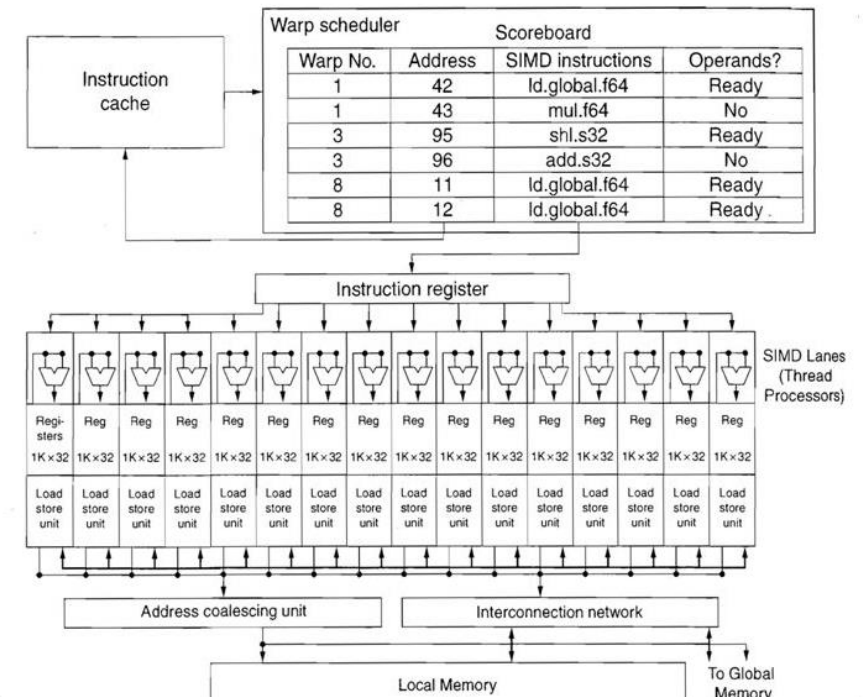


SIMD REGISTERS

Their approach has been multi-pronged:



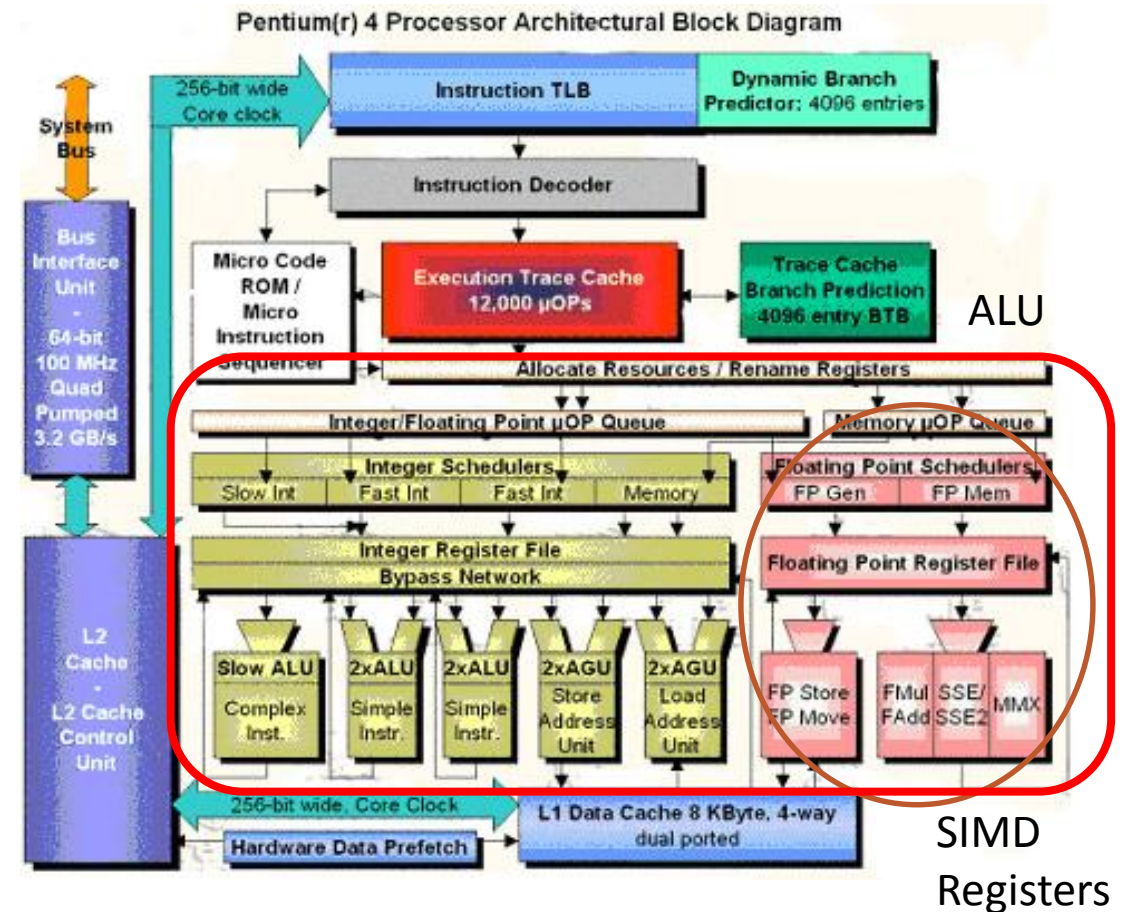
Each CPU contains numerous cores..



..and each core contains SIMD registers.

MMX


The first SIMD registers were introduced as by Intel in 1997.



MMX Concept

The concept – that data would be packed together into a single MMX type.

A single operation on an MMX type would correspond to that operation being performed on the data packed within in.



MMX Concept

Explained in another way:

Scalar Instructions

$$\begin{array}{rcl} 4 & + & 1 = 5 \\ 0 & + & 3 = 3 \\ -2 & + & 8 = 6 \\ 9 & + & -7 = 2 \end{array}$$

4 separate operations
(additions) on integers

Vector Instructions

$$\begin{array}{rcl} 4 & 1 & 5 \\ 0 & 3 & 3 \\ -2 & 8 & 6 \\ 9 & -7 & 2 \end{array} \quad + \quad \begin{array}{rcl} 1 & 3 & 8 & -7 \end{array} = \begin{array}{rcl} 5 & 3 & 6 & 2 \end{array}$$

Vector Length

1 single operation (addition)
performed on packed data.

MMX Packed Vector
Length = 64 bits.



From MMX to SSE

These were very successful – however, 64 bits was too small, and MMX was only useful for integer computations.

These parallel registers were upgraded to:

- Hold more data (128 bits), and
- Perform floating point operations.

and SSE was born.



SSE Intrinsic Functions

SSE has been incredibly successful – so successful, unfortunately, that most compilers default to it automatically, even nowadays.

We can get information on the SSE functions from the Intel Intrinsic Guide.

SSE has since been superseded by AVX.



Technologies

☐ MMX☐ SSE☐ SSE2☐ SSE3☐ SSSE3☐ SSE4.1☐ SSE4.2☐ AVX☐ AVX2☐ FMA☐ AVX-512☐ KNC☐ SVML☐ Other

Categories

☐ Application-Targeted

The Intel Intrinsic Guide is an interactive reference tool for many Intel instructions - including Intel® SSE, AVX, AVX-

```
int _mm512_4dpwssd_epi32 (_m512i src, _m512i mask)
```

```
int _mm512_mask_4dpwssd_epi32 (_m512i src, _mmask16 mask, _m512i a3, _m128i * b)
```

```
int _mm512_maskz_4dpwssd_epi32 (_mmask16 mask, _m512i a3, _m128i * b)
```

```
int _mm512_4dpwssds_epi32 (_m512i src, _m512i mask)
```

```
int _mm512_mask_4dpwssds_epi32 (_m512i src, _mmask16 mask, _m512i a3, _m128i * b)
```

```
int _mm512_maskz_4dpwssds_epi32 (_m512i src, _mmask16 mask, _m512i a3, _m128i * b)
```

```
__m512 _mm512_4fmadd_ps (_m512 a, _m512i b, _m512i c)
```

```
__m512 _mm512_mask_4fmadd_ps (_m512 a, _mmask16 mask, _m512 b, _m128i * c)
```

```
__m512 _mm512_maskz_4fmadd_ps (_m512 a, _mmask16 mask, _m512 b, _m128i * c)
```

Advanced Vector Extensions

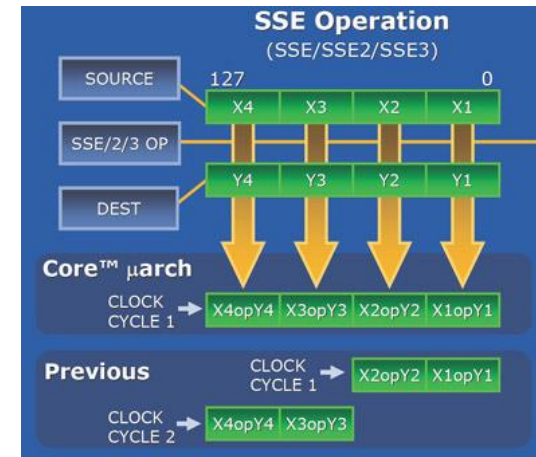
AVX stands for Advanced Vector eXtensions.

It is basically an extension of the SSE instruction set – the concept is the same.

However, AVX registers contain 256 bits of data.



MMX - 64 bit wide registers for parallel computing on integer types.



SSE - 128 bit wide registers for parallel computing on integers, floats and (later) doubles.

The advertisement features the headline "GET TO THE POINT" and "NEW SERVERS BASED ON THE INTEL® XEON® PROCESSOR E5-2600 PRODUCT FAMILY". It includes the Intel Xeon logo. Below the headline, there is a graphic of a person pointing at a whiteboard with a mathematical formula, and another person sitting at a desk with a computer. The text "NO INTEL® AVX" and "INTEL® AVX" is shown. The main claim is "GET UP TO 2X FASTER WITH INTEL® ADVANCED VECTOR EXTENSIONS (INTEL® AVX)". At the bottom, there is a small disclaimer: "For complete configuration and performance details, refer to http://www.intel.com/xeon. Intel, the Intel logo, Xeon, and Xeon Inside are trademarks of Intel Corporation in the U.S. and/or other countries. Other names and brands may be claimed as the property of others. Copyright © 2012 Intel Corporation. All rights reserved."

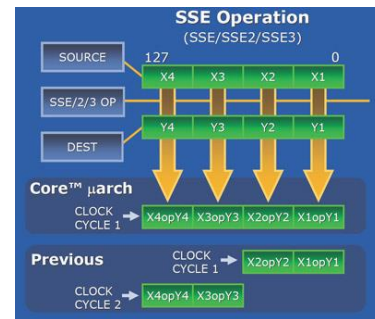
AVX - 256 bit wide registers for parallel computing on integers, floats and doubles.

AVX-512 register scheme as extension from the AVX (YMM0-YMM15) and SSE (XMM0-XMM15) registers

511	256	255	128	127	0
ZMM0	YMM0	XMM0			
ZMM1	YMM1	XMM1			
ZMM2	YMM2	XMM2			
ZMM3	YMM3	XMM3			
ZMM4	YMM4	XMM4			
ZMM5	YMM5	XMM5			
ZMM6	YMM6	XMM6			
ZMM7	YMM7	XMM7			
ZMM8	YMM8	XMM8			
ZMM9	YMM9	XMM9			
ZMM10	YMM10	XMM10			
ZMM11	YMM11	XMM11			
ZMM12	YMM12	XMM12			
ZMM13	YMM13	XMM13			
ZMM14	YMM14	XMM14			
ZMM15	YMM15	XMM15			
ZMM16	YMM16	XMM16			

Quick Guide

XMM



SSE - 128 bit wide registers for parallel computing on integers, floats and (later) doubles.

YMM

GET TO THE POINT

NEW SERVERS BASED ON THE INTEL® XEON® PROCESSOR E5-2600 PRODUCT FAMILY

NO INTEL® AVX INTEL® AVX

GET UP TO **2X FASTER** WITH INTEL® ADVANCED VECTOR EXTENSIONS (INTEL® AVX)¹

For complete configuration and performance details, refer to <http://www.intel.com/content/www/us/en/processors/xeon/xeon-processor-e5-2600-product-family.html>.
¹ Other names and brands may be claimed as the property of others. Copyright © 2012 Intel Corporation. All rights reserved.

AVX - 256 bit wide registers for parallel computing on integers, floats and doubles.

ZMM

INTEL® ADVANCED VECTOR EXTENSIONS 512

IN INTEL® XEON® SCALABLE PROCESSORS

Broadwell — Knights Landing — Skylake


Comparative analysis of vector processing functionality

AVX 512 - 512 bit wide registers for parallel computing on integers, floats and doubles.


Advanced Vector Extensions

The additional space in the registers means that, instead of performing simultaneous operations across **4 floats**, we can now do the same for **8 floats**.


This means that we should see a vast improvement in performance!



One single AVX Addition



One single AVX assignment



One single AVX packed result

Getting Started

Several things need to happen before we can use AVX to accelerate our computations:

- We need to find a better way to allocate memory.
- We need to pack our arrays of floats into arrays of AVX types.
- We need to use AVX intrinsic functions for our calculation.

Piece of cake. Let's deal with memory allocation first.

Memory Allocation

The code we started with used malloc:

```
// Allocate memory  
a = (float*)malloc(size); b = (float*)malloc(size); c = (float*)malloc(size);
```

Malloc will align our data in memory based on the type of data being allocated.

However, if we want to pack our data into larger types, we need align our data along cache boundaries based on the larger type.

Memory Allocation

Being old-school, I use posix_memalign for this:

```
// Allocate memory along cache boundaries
posix_memalign((void**)&a, alignment, size);
posix_memalign((void**)&b, alignment, size);
posix_memalign((void**)&c, alignment, size);
```

where `size_t size = 32*sizeof(float)` (32 byte alignment) – this will be OK for AVX, but we might need to increase it again if the vector length changes. (HINT)

NOTE: `posix_memalign` is found inside `<stdlib.h>`, and the variables are freed in the same way [`free(a)` etc].

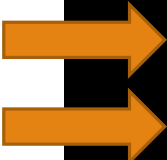
AVX Packing and Computation

Let's create a new function – `Compute_C_AVX()` – to perform the same job as our previous code.

First things first, we need to pack our floats into 256 bit wide AVX types – the fastest way to do this is to perform a direct type conversion / casting using an AVX pointer:

`__m256 *AVX_a;` `// Pointer to AVX type`

`AVX_a = (__m256*)x;` `// Type conversion (casting)`



```
void Compute_C_AVX(float *x, float *y, float *z, int N) {  
  
    int i;  
    int N_AVX = N/8; // Number of AVX packed data types. Should be a multiple of 8  
    __m256 *AVX_a, *AVX_b, *AVX_c; // 256 bit packed type arrays  
    __m256 AVX_tmp;                // 256 bit single tmp value  
    // Map x, y, z onto these  
    AVX_a = (__m256*)x; AVX_b = (__m256*)y; AVX_c = (__m256*)z;  
    // Compute result for each packed type  
    for (i = 0; i < N_AVX; i++) {  
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];  
        AVX_c[i] = _mm256_sqrt_ps(AVX_tmp);  
    }  
}  
  
void Compute_C(float *x, float *y, float *z, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        z[i] = sqrtf(x[i]*x[i] + y[i]*y[i]);  
    }  
}
```


AVX Packing and Computation

Each time we perform a computation on an AVX type, we are performing a computation on 8 packed floats.

Since our previous problem size was 1024, this means we will need to iterate over $1024/8 = 128$ AVX types.

Our for loop needs to reflect this:

```
N_AVX = N/8;
for (i = 0; i < N_AVX; i++) {
    <stuff>
}
```



```
void Compute_C_AVX(float *x, float *y, float *z, int N) {
    int i;
    int N_AVX = N/8; // Number of AVX packed data types. Should be a
    __m256 *AVX_a, *AVX_b, *AVX_c; // 256 bit packed type arrays
    __m256 AVX_tmp; // 256 bit single tmp value
    // Map x, y, z onto these
    AVX_a = (__m256*)x; AVX_b = (__m256*)y; AVX_c = (__m256*)z;
    // Compute result for each packed type
    for (i = 0; i < N_AVX; i++) {
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];
        AVX_c[i] = _mm256_sqrt_ps(AVX_tmp);
    }
}

void Compute_C(float *x, float *y, float *z, int N) {
    int i;
    for (i = 0; i < N; i++) {
        z[i] = sqrtf(x[i]*x[i] + y[i]*y[i]);
    }
}
```

AVX Packing and Computation

Finally, the computation.


The squaring and addition is straightforward:

`AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];`

AVX_tmp contains 8 packed floats now. We cannot use the ordinary sqrt function – we need a function to use on an AVX type:

`AVX_c[i] = _mm256_sqrt_ps(AVX_tmp);`

`_mm256_sqrt_ps` is a vector intrinsic function



```
void Compute_C_AVX(float *x, float *y, float *z, int N) {  
  
    int i;  
    int N_AVX = N/8; // Number of AVX packed data types. Should be a  
    __m256 *AVX_a, *AVX_b, *AVX_c; // 256 bit packed type arrays  
    __m256 AVX_tmp; // 256 bit single tmp value  
    // Map x, y, z onto these  
    AVX_a = (__m256*)x; AVX_b = (__m256*)y; AVX_c = (__m256*)z;  
    // Compute result for each packed type  
    for (i = 0; i < N_AVX; i++) {  
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];  
        AVX_c[i] = _mm256_sqrt_ps(AVX_tmp);  
    }  
}  
  
void Compute_C(float *x, float *y, float *z, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        z[i] = sqrtf(x[i]*x[i] + y[i]*y[i]);  
    }  
}
```


AVX Packing and Computation

Quick review of our new function using AVX types and intrinsic functions:

- Our loop goes from 1 to N_AVX, performing 1/8th the iterations than normal since we are performing operations on 8 packed floats each iteration.
- We use the special vector intrinsic function `_mm256_sqrt_ps` to compute the sqrt on each packed float in `AVX_tmp`.

```
void Compute_C_AVX(float *x, float *y, float *z, int N) {  
  
    int i;  
    int N_AVX = N/8; // Number of AVX packed data types. Should be a  
    __m256 *AVX_a, *AVX_b, *AVX_c; // 256 bit packed type arrays  
    __m256 AVX_tmp; // 256 bit single tmp value  
    // Map x, y, z onto these  
    AVX_a = (__m256*)x; AVX_b = (__m256*)y; AVX_c = (__m256*)z;  
    // Compute result for each packed type  
    for (i = 0; i < N_AVX; i++) {  
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];  
        AVX_c[i] = _mm256_sqrt_ps(AVX_tmp);  
    }  
}  
  
void Compute_C(float *x, float *y, float *z, int N) {  
    int i;  
    for (i = 0; i < N; i++) {  
        z[i] = sqrtf(x[i]*x[i] + y[i]*y[i]);  
    }  
}
```

AVX Packing and Computation

To make this code, we need to modify our makefile by adding a flag:

```
all:
    g++ main.cpp -O3 -mavx -c
    g++ main.o -O3 -mavx -o test.run
```

Let's rebuild and see what happens.

AVX Code & Result

Perform a quick object dump → let's see what the computer is doing now:

- We see the computer is now using the larger (256 bit) YMM0 and YMM1 registers.
- We also see that the instructions we are performing are all parallel instructions:

(vmulps, vaddps, vsqrtps etc)

This tells us we are making proper use of AVX vectorization.

```
0000000000000000 <_Z13Compute_C_AVXPfS_S_i>:
 0: 8d 41 07          lea     0x7(%rcx),%eax
 3: 85 c9             test    %ecx,%ecx
 5: 0f 48 c8          cmovs   %eax,%ecx
 8: c1 f9 03          sar     $0x3,%ecx
 b: 85 c9             test    %ecx,%ecx
 d: 7e 3c             jle     4b <_Z13Compute_C_AVXPfS_S_i+0x4b>
 f: 83 e9 01          sub     $0x1,%ecx
12: 31 c0             xor     %eax,%eax
14: 48 83 c1 01       add     $0x1,%rcx
18: 48 c1 e1 05       shl     $0x5,%rcx
1c: 0f 1f 40 00       nopl    0x0(%rax)
20: c5 fc 28 0c 07    vmovaps (%rdi,%rax,1),%ymm1
25: c5 fc 28 04 06    vmovaps (%rsi,%rax,1),%ymm0
2a: c5 f4 59 c9       vmulps  %ymm1,%ymm1,%ymm1
2e: c5 fc 59 c0       vmulps  %ymm0,%ymm0,%ymm0
32: c5 f4 58 c0       vaddps  %ymm0,%ymm1,%ymm0
36: c5 fc 51 c0       vsqrtps %ymm0,%ymm0
3a: c5 fc 29 04 02    vmovaps %ymm0, (%rdx,%rax,1)
3f: 48 83 c0 20       add     $0x20,%rax
43: 48 39 c8          cmp     %rcx,%rax
46: 75 d8             jne     20 <_Z13Compute_C_AVXPfS_S_i+0x20>
48: c5 f8 77         vzeroupper
4b: f3 c3             repz retq
4d: 0f 1f 00         nopl    (%rax)
```

AVX Code & Result

Compare the two assembly-like codes:

```
0000000000000000 <_Z13Compute_C_AVXPfS_S_i>:
 0: 8d 41 07          lea     0x7(%rcx),%eax
 3: 85 c9             test    %ecx,%ecx
 5: 0f 48 c8          cmovs   %eax,%ecx
 8: c1 f9 03          sar     $0x3,%ecx
 b: 85 c9             test    %ecx,%ecx
 d: 7e 3c             jle     4b <_Z13Compute_C_AVXPfS_S_i+0x4b>
 f: 83 e9 01          sub     $0x1,%ecx
12: 31 c0             xor     %eax,%eax
14: 48 83 c1 01        add     $0x1,%rcx
18: 48 c1 e1 05        shl     $0x5,%rcx
1c: 0f 1f 40 00        nopl    0x0(%rax)
20: c5 fc 28 0c 07     vmovaps (%rdi,%rax,1),%ymm1
25: c5 fc 28 04 06     vmovaps (%rsi,%rax,1),%ymm0
2a: c5 f4 59 c9        vmulps  %ymm1,%ymm1,%ymm1
2e: c5 fc 59 c0        vmulps  %ymm0,%ymm0,%ymm0
32: c5 f4 58 c0        vaddps  %ymm0,%ymm1,%ymm0
36: c5 fc 51 c0        vsqrtps %ymm0,%ymm0
3a: c5 fc 29 04 02     vmovaps %ymm0, (%rdx,%rax,1)
3f: 48 83 c0 20        add     $0x20,%rax
43: 48 39 c8           cmp     %rcx,%rax
46: 75 d8             jne     20 <_Z13Compute_C_AVXPfS_S_i+0x20>
48: c5 f8 77          vzeroupper
4b: f3 c3             repz    retq
4d: 0f 1f 00          nopl    (%rax)
```

avx function

```
0000000000000050 <_Z9Compute_CPfS_S_i>:
50: 85 c9             test    %ecx,%ecx
52: 7e 50             jle     a4 <_Z9Compute_CPfS_S_i+0x54>
54: 55               push    %rbp
55: 8d 41 ff          lea     -0x1(%rcx),%eax
58: 53               push    %rbx
59: 48 8d 2c 85 04 00 00 lea     0x4(%rax,4),%rbp
60: 00
61: 31 db            xor     %ebx,%ebx
63: 48 83 ec 28       sub     $0x28,%rsp
67: 66 0f 1f 84 00 00 00 nopw    0x0(%rax,%rax,1)
6e: 00 00
70: c5 fa 10 0c 1f     vmovss  (%rdi,%rbx,1),%xmm1
75: c5 fa 10 04 1e     vmovss  (%rsi,%rbx,1),%xmm0
7a: c5 f2 59 c9        vmulss  %xmm1,%xmm1,%xmm1
7e: c5 fa 59 c0        vmulss  %xmm0,%xmm0,%xmm0
82: c5 f2 58 c0        vaddss  %xmm0,%xmm1,%xmm0
86: c5 f2 51 c8        vsqrts  %xmm0,%xmm1,%xmm1
8a: c5 f8 2e c9        vucomiss %xmm1,%xmm1
8e: 7a 16            jp      a6 <_Z9Compute_CPfS_S_i+0x56>
90: c5 fa 11 0c 1a     vmovss  %xmm1, (%rdx,%rbx,1)
95: 48 83 c3 04        add     $0x4,%rbx
99: 48 39 eb          cmp     %rbp,%rbx
9c: 75 d2             jne     70 <_Z9Compute_CPfS_S_i+0x20>
9e: 48 83 c4 28        add     $0x28,%rsp
a2: 5b               pop     %rbx
a3: 5d               pop     %rbp
a4: f3 c3            repz    retq
a6: 48 89 54 24 18     mov     %rdx,0x18(%rsp)
ab: 48 89 74 24 10     mov     %rsi,0x10(%rsp)
b0: 48 89 7c 24 08     mov     %rdi,0x8(%rsp)
b5: e8 00 00 00 00     callq   ba <_Z9Compute_CPfS_S_i+0x6a>
ba: 48 8b 54 24 18     mov     0x18(%rsp),%rdx
bf: c5 f8 28 c8        vmovaps %xmm0,%xmm1
c3: 48 8b 74 24 10     mov     0x10(%rsp),%rsi
c8: 48 8b 7c 24 08     mov     0x8(%rsp),%rdi
cd: eb c1            jmp     90 <_Z9Compute_CPfS_S_i+0x40>
```

normal
function

Whole Code

```
/* Sequential demonstration prepared for AVX Cookies 'n' Code Session
   Dr. Matthew Smith, msmith@astro.swin.edu.au */

#include <stdio.h>
#include <stdlib.h> // For posix_memalign
#include <math.h>
#include <immintrin.h> // For AVX Intrinsic functions and types

void Compute_C(float *x, float *y, float *z, int N);
void Compute_C_AVX(float *x, float *y, float *z, int N);

int main() {
    int N = 1024; // N = size of problem, a power of 2 is good.
    float *a, *b, *c; // Three vectors containing N elements
    float error = 0.0; // Used later for testing
    size_t size = N*sizeof(float);
    size_t alignment = 32;
    int i;
    // Allocate memory along cache boundaries
    posix_memalign((void**)&a, alignment, size);
    posix_memalign((void**)&b, alignment, size);
    posix_memalign((void**)&c, alignment, size);

    // Initialise a and b vectors
    for (i = 0; i < N; i++) {
        a[i] = (float)i; b[i] = (float)(2*i);
    }

    Compute_C_AVX(a, b, c, N);

    // Check the result
    for (i = 0; i < N; i++) {
        error = error + c[i] - sqrtf(a[i]*a[i] + b[i]*b[i]);
    }
    printf("Error = %g\n", error);

    free(a); free(b); free(c);
    printf("Computation Complete\n");
    return 0;
}
```

```
void Compute_C_AVX(float *x, float *y, float *z, int N) {

    int i;
    int N_AVX = N/8; // Number of AVX packed data types. Should be a power of 2
    __m256 *AVX_a, *AVX_b, *AVX_c; // 256 bit packed type arrays
    __m256 AVX_tmp; // 256 bit single tmp value
    // Map x, y, z onto these
    AVX_a = (__m256*)x; AVX_b = (__m256*)y; AVX_c = (__m256*)z;
    // Compute result for each packed type
    for (i = 0; i < N_AVX; i++) {
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];
        AVX_c[i] = _mm256_sqrt_ps(AVX_tmp);
    }
}

void Compute_C(float *x, float *y, float *z, int N) {
    int i;
    for (i = 0; i < N; i++) {
        z[i] = sqrtf(x[i]*x[i] + y[i]*y[i]);
    }
}
```

AVX512

AVX512 on Ozstar

In Ozstar we are fortunate enough to be using a skylake architecture.

This means that the Intel Gold 6140 supports AVX512 – meaning the registers inside its ALU have a length of 512 bits!

Intel® TSX-NI ?	Yes
Intel® 64 ?	Yes
Instruction Set Extensions ?	Intel® SSE4.2, Intel® AVX, Intel® AVX2, Intel® AVX-512
# of AVX-512 FMA Units ?	2
Enhanced Intel SpeedStep® Technology ?	Yes

OzSTAR specifications

107 Standard compute nodes:

Dell R740 14G Server

2 x Intel Gold 6140 18-core processors

AVX-512 Code

Altering code from AVX to AVX-512 is pretty easy (easier than from SSE to AVX):

- Our `N_AVX` variable is now `N/16` since we pack 16 singles into 512 bits.
- Where we once used 256, we now use 512. Find and replace job, its easy enough.

We also need to make sure our memory is aligned to 64 bytes and **not** 32.

```
void Compute_C_AVX512(float *x, float *y, float *z, int N) {  
  
    int i;  
    int N_AVX = N/16; // Number of AVX512 packed data types. Should be  
    __m512 *AVX_a, *AVX_b, *AVX_c; // 512 bit packed type arrays  
    __m512 AVX_tmp;                // 512 bit single tmp value  
    // Map x, y, z onto these  
    AVX_a = (__m512*)x; AVX_b = (__m512*)y; AVX_c = (__m512*)z;  
    // Compute result for each packed type  
    for (i = 0; i < N_AVX; i++) {  
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];  
        AVX_c[i] = _mm512_sqrt_ps(AVX_tmp);  
    }  
}
```

```
size_t alignment = 64;  
int i;  
// Allocate memory along cache boundaries  
posix_memalign((void**)&a, alignment, size);  
posix_memalign((void**)&b, alignment, size);  
posix_memalign((void**)&c, alignment, size);
```

Requirements

The version of g++ which is loaded by default when you log into Ozstar won't cut it for AVX512.

You need to load the newer version: **module load gcc/7.3.0**

We also need to update our makefile to let g++ know we wanna use skylake's AVX 512:

```
all:
    g++ main.cpp -O3 -mavx -march=skylake-avx512 -c
    g++ main.o -O3 -mavx -march=skylake-avx512 -o test.run
```

Let's rebuild it and see what happens now.

AVX512 Review

Let's examine the object disassembly:

- The AVX512 code is using the ZMM0 and ZMM1 registers.
- All operations are still being performed in parallel.

Looks like we nailed it.

This code will outperform any attempt at auto-vectorization the compiler will attempt, especially for complex computations (not covered here).

```
0000000000000000 <_Z13Compute_C_AVXPfS_S_i>:
  0: 85 c9          test    %ecx,%ecx
  2: 8d 41 07       lea     0x7(%rcx),%eax
  5: 0f 48 c8       cmovs  %eax,%ecx
  8: c1 f9 03       sar     $0x3,%ecx
  b: 85 c9          test    %ecx,%ecx
  d: 7e 39         jle     48 <_Z13Compute_C_AVXPfS_S_i+0x48>
  f: 83 e9 01       sub     $0x1,%ecx
 12: 31 c0          xor     %eax,%eax
 14: 48 83 c1 01    add     $0x1,%rcx
 18: 48 c1 e1 05    shl     $0x5,%rcx
 1c: 0f 1f 40 00    nopl   0x0(%rax)
 20: c5 fc 28 0c 06 vmovaps (%rsi,%rax,1),%ymm1
 25: c5 fc 28 04 07 vmovaps (%rdi,%rax,1),%ymm0
 2a: c5 f4 59 c9    vmulps %ymm1,%ymm1,%ymm1
 2e: c4 e2 75 98 c0 vfmadd132ps %ymm0,%ymm1,%ymm0
 33: c5 fc 51 c0    vsqrtps %ymm0,%ymm0
 37: c5 fc 29 04 02 vmovaps %ymm0, (%rdx,%rax,1)
 3c: 48 83 c0 20    add     $0x20,%rax
 40: 48 39 c1       cmp     %rax,%rcx
 43: 75 db         jne     20 <_Z13Compute_C_AVXPfS_S_i+0x20>
 45: c5 f8 77       vzeroupper
 48: c3            retq
 49: 0f 1f 80 00 00 00 00 nopl   0x0(%rax)

0000000000000050 <_Z16Compute_C_AVX512PfS_S_i>:
 50: 85 c9          test    %ecx,%ecx
 52: 8d 41 0f       lea     0xf(%rcx),%eax
 55: 0f 48 c8       cmovs  %eax,%ecx
 58: c1 f9 04       sar     $0x4,%ecx
 5b: 85 c9          test    %ecx,%ecx
 5d: 7e 44         jle     a3 <_Z16Compute_C_AVX512PfS_S_i+0x53>
 5f: 83 e9 01       sub     $0x1,%ecx
 62: 31 c0          xor     %eax,%eax
 64: 48 83 c1 01    add     $0x1,%rcx
 68: 48 c1 e1 06    shl     $0x6,%rcx
 6c: 0f 1f 40 00    nopl   0x0(%rax)
 70: 62 f1 7c 48 28 0c 06 vmovaps (%rsi,%rax,1),%zmm1
 77: 62 f1 7c 48 28 04 07 vmovaps (%rdi,%rax,1),%zmm0
 7e: 62 f1 74 48 59 c9    vmulps %zmm1,%zmm1,%zmm1
 84: 62 f2 75 48 98 c0    vfmadd132ps %zmm0,%zmm1,%zmm0
 8a: 62 f1 7c 48 51 c0    vsqrtps %zmm0,%zmm0
 90: 62 f1 7c 48 29 04 02 vmovaps %zmm0, (%rdx,%rax,1)
 97: 48 83 c0 40    add     $0x40,%rax
 9b: 48 39 c1       cmp     %rax,%rcx
 9e: 75 d0         jne     70 <_Z16Compute_C_AVX512PfS_S_i+0x20>
 a0: c5 f8 77       vzeroupper
 a3: c3            retq
 a4: 66 90         xchg    %ax,%ax
 a6: 66 2e 0f 1f 84 00 00 nopw    %cs:0x0(%rax,%rax,1)
 ad: 00 00 00
```

Padding Length

One problem which has loomed over our heads is the requirement that N – our problem size – be evenly divisible by 16 (when using floats).

The solution is to pad the memory length – allocate more memory than you need so that the memory you do have is a multiple of 16 in length.

The computation is quite trivial: $NT = 16 \left[(int) \left(\frac{N - 16 + 1}{16} \right) \right]$

Example: if $N = 1025$, then NT will be 1040 (1024+16) → giving us $15 * 4 = 60$ wasted bytes, but no segmentation error.

Final Product

Final code using:

- Cache boundary aligned memory
- Padded memory up to multiple of 16
- AVX512 packed types
- AVX512 Intrinsic Functions

This code will run much faster than the code we started out with.

```
void Compute_C_AVX512(float *x, float *y, float *z, int N) {
    int i;
    int N_AVX = N/16; // Integer number of AVX512 packed types.
    __m512 *AVX_a, *AVX_b, *AVX_c; // 512 bit packed type arrays
    __m512 AVX_tmp; // 512 bit single tmp value
    // Map x, y, z onto these
    AVX_a = (__m512*)x; AVX_b = (__m512*)y; AVX_c = (__m512*)z;
    // Compute result for each packed type
    for (i = 0; i < N_AVX; i++) {
        AVX_tmp = AVX_a[i]*AVX_a[i] + AVX_b[i]*AVX_b[i];
        AVX_c[i] = _mm512_sqrt_ps(AVX_tmp);
    }
}
```

```
#include <stdio.h>
#include <stdlib.h> // For posix_memalign
#include <math.h>
#include <immintrin.h> // For AVX Intrinsic functions and types

void Compute_C_AVX512(float *x, float *y, float *z, int N);

int main() {
    int N = 1025; // N = size of problem, a power of 2 is good.
    int NT;
    float *a, *b, *c; // Three vectors containing N elements
    float error = 0.0; // Used later for testing
    size_t size;
    size_t alignment = 64;
    int i;

    // Compute Padded Length
    NT = (int) ((N+16-1)/16);
    NT = 16*NT;
    size = NT*sizeof(float);
    // Print for our confirmation
    printf("Original problem size = %d, Modified problem size = %d\n", N, NT);

    // Allocate memory along cache boundaries
    posix_memalign((void**)&a, alignment, size);
    posix_memalign((void**)&b, alignment, size);
    posix_memalign((void**)&c, alignment, size);

    // Initialise a and b vectors
    for (i = 0; i < N; i++) {
        a[i] = (float)i; b[i] = (float)(2*i);
    }

    // Use the padded length when calling our AVX512 function
    Compute_C_AVX512(a, b, c, NT);

    // Check the result
    for (i = 0; i < N; i++) {
        error = error + c[i] - sqrtf(a[i]*a[i] + b[i]*b[i]);
    }
    printf("Error = %g\n", error);

    free(a); free(b); free(c);
    printf("Computation Complete\n");
    return 0;
}
```

Homework

Check out the following possibilities:

- Rewrite the example to employ OpenMP parallelization in addition to AVX vectorization.
- Repeat the above step with MPI.

And that's it. Hopefully by now I have had enough cookies to tide me over until next time.

```
#include <stdio.h>
#include <omp.h>
#include <immintrin.h>

#define P 4
#define N 32

float *a;

int main() {
    size_t alignment = 32;
    int error = posix_memalign((void**)&a, alignment, N*sizeof(float));

    for (int i = 0; i < N; i++) {
        a[i] = (float)i;
    }

    omp_set_num_threads(P);

    #pragma omp parallel
    {
        int tid = omp_get_thread_num();
        int index = tid;

        printf("Index = %d\n", index);
        // Declare AVX_a here
        __m256 *AVX_a;
        // Set AVX_a
        AVX_a = (__m256*)a + index;
        __m256 AVX_b = _mm256_set1_ps(2.0);

        AVX_a[0] = _mm256_mul_ps(AVX_a[0], AVX_b);
    }

    for (int i = 0; i < N; i++) {
        printf("a[%d] = %g\n", i, a[i]);
    }

    free(a);
}
```